

eda_19-02-2021

February 17, 2021

```
[1]: import sys, os
sys.path.append("C:/users/kristijan/documents/projects/jhu1/")
sys.path.append("C:/users/kristijan/documents/projects/jhu1/data")
sys.path.append("C:/users/kristijan/documents/projects/jhu1/output")
os.chdir("C:/users/kristijan/documents/projects/jhu1/")
import pandas as pd
import numpy as np
import plotly.express as px
import seaborn as sns
sns.set(font_scale=1.3)
from matplotlib import pyplot as plt
from stimuli import prefixes, prompts
import json
```

1 Introduction

We are studying the functional properties of the memory systems allowing ANNs (Transformers and RNNs) to solve natural language tasks.

Cognitive neuroscience provides us with two possible scenarios:

- A) memory traces are **short lived** akin to the **working memory system** posited by Baddeley et al (i.e. isolated from long-term knowledge)
- B) memory traces are **long-lived** and entangled/or interacting with long-term knowledge (e.g. **hippocampal/episodic system**)

We study these dimensions of a working memory task:

- **list length** (nr. of items to be remembered)
- **prompt length** (prompt) length: number of items intervening between the remembered list and repeated exposure
- **list composition** (list type): whether nouns in the list share semantic category (categorized or random)
- **context (prompt) type** whether or not prompt is intact or scrambled (incongruent)

Insofar as predictions are concerned, it seems the A) and B) should primarily differ in terms of the influence of semantic knowledge on performance:

- A): shared semantic properties (of word lists, of prompts) lead to **improved performance** as short-lived memory traces are **refreshed** (like rehearsal)
- B): shared semantic properties (of word lists, of prompts) lead to **impaired performance** as similarity in long-lasting traces leads to **interference** (too many similar traces active)

(Not so sure about list length and other dimensions atm).

1.1 Loading in the data and some minimal wrangling

```
[2]: # read in the outputs created in outputs4dataframe.py
gpt = pd.read_csv("./output/output_gpt2.csv", sep="\t")
rnn = pd.read_csv("./output/output_rnn.csv", sep="\t")

# rename prompt length values to more meaningful ones
prompt_len_map = {
    1:8,
    2:30,
    3:100,
    4:200,
    5:400,
}
gpt.prompt_len = gpt.prompt_len.map(prompt_len_map)
rnn.prompt_len = rnn.prompt_len.map(prompt_len_map)

# rename scenario values to more meaningful ones
scenario_map = {
    "sce1":"intact",
    "sce1rnd": "scrambled",
}
gpt.scenario = gpt.scenario.map(scenario_map)
rnn.scenario = rnn.scenario.map(scenario_map)

# rename the "scenario" column to "context"
gpt.rename(columns={"scenario": "context"}, inplace=True)
rnn.rename(columns={"scenario": "context"}, inplace=True)

# add token position offsets, to start indexing with 1
gpt.marker_pos += 1
rnn.marker_pos += 1

# drop some redundant columns created by Pandas bookkeeping system
rnn.drop(["Unnamed: 0"], axis=1, inplace=True);
gpt.drop(["Unnamed: 0"], axis=1, inplace=True);
gpt.drop(["Unnamed: 0.1"], axis=1, inplace=True);
```

1.2 Data checks

First, we check that all seems good on the input data side.

We will be interested in these variables:

- surp: surprisal
- list_len: length of the list
- prompt_len: length of the prompt between two lists (n. tokens)

- `list`: list composition (categorized or random)
- `second_list`: whether the second list is repeated, permuted or an unrelated (control) list

RNN output (punctuation is removed):

```
[3]: rnn.head(5)
```

```
[3]:
```

	word	sentid	corpuspos	marker	prompt_len	list_len	wlen	surp	\
0	before	0	0	0	8	3.0	6	10.156857	
1	the	0	1	0	8	3.0	3	1.502160	
2	meeting	0	2	0	8	3.0	7	8.726244	
3	<unk>	0	4	0	8	3.0	5	1.729451	
4	wrote	0	5	0	8	3.0	5	6.610217	

	hs	dHs	list	second_list	context	marker_pos
0	6.533242	0.000000	categorized	permute	intact	1
1	5.209762	1.323480	categorized	permute	intact	2
2	7.835398	-2.625636	categorized	permute	intact	3
3	4.781773	0.529446	categorized	permute	intact	4
4	7.430719	-2.648947	categorized	permute	intact	5

GPT-2 output:

```
[4]: gpt.head(5)
```

```
[4]:
```

	token	marker	positionID	surp	ispunct	prefix	prompt_len	\
0	Before	0	0	NaN	False	1	8	
1	the	0	1	2.551027	False	1	8	
2	meeting	0	2	5.753714	False	1	8	
3	Mary	0	4	8.620794	False	1	8	
4	wrote	0	5	6.034746	False	1	8	

	list_len	sentid	second_list	list	context	sentpos	marker_pos
0	3	1	permute	categorized	scrambled	0	1
1	3	1	permute	categorized	scrambled	1	2
2	3	1	permute	categorized	scrambled	2	3
3	3	1	permute	categorized	scrambled	3	4
4	3	1	permute	categorized	scrambled	4	5

1.3 Checking the stimuli and prompt manipulation

First, let's check the input strings for the congruent scenario:

```
[5]: stimulus_id = 92 # let's pick a trial
selection = ((gpt.sentid==stimulus_id) & (gpt.second_list=="permute") & (gpt.
    ↳list=="random") & (gpt.context=="intact"))
for i in range(4):
    print("{}: ".format(i) + " ".join(gpt.loc[(gpt.marker == i) & selection].
    ↳token.tolist()).strip())
```

0: Before the meeting Mary wrote down the following list of words
 1: farmer contest outline success orange
 2: After the meeting Mary went for a walk It was a busy day and she needed a break Outside was really beautiful and warm and the flowers in the park were blo oming When she got back she read the list again
 3: success outline contest farmer orange

So the prompt under 2 looks fine.

We now do the same but we select rows marked as incongruent prompt and string under 2 should be permuted. The two lists are the same, only the second one is permuted.

```
[6]: selection = ((gpt.sentid==stimulus_id) & (gpt.second_list=="permute") & (gpt.
    ↪list=="random") & (gpt.context=="scrambled"))
for i in range(4):
    print("{}: ".format(i) + " ".join(gpt.loc[(gpt.marker == i) & selection].
    ↪token.tolist()).strip())
```

0: Before the meeting Mary wrote down the following list of words
 1: farmer contest outline success orange
 2: the back in for and were a When she walk was meeting Outside break the got warm went and After Mary the read list park beautiful again really a a needed It busy was she and she day the flowers blo oming
 3: success outline contest farmer orange

2 Results

2.1 RNN

We now explore how surprisal fluctuates over token positions (averaged over 20 word lists, per position).

We put token position on x-axes and surprisal on y.

We color code prompt length.

Prompt type – congruent or incongruent – is show in top and bottom rows, respectively.

List type – repeated, permuted or control – is displayed from left to right column-wise.

Points mark average surprisal over 20 tokens, bars denote 95% confidence intervals (bootstrapped by sampling with replacement 1000 times).

We start with RNN. We make 3 plots (list lengths of 10, 5, and 3).

We see there is little variability going from left to right column-wise meaning that in this paradigm we cannot detect RNN sensitivity to list identity or word order.

On the other hand, we do see that bottom row has a sort of a surprisal jump for the lightest color (max prompt length of 400 tokens).

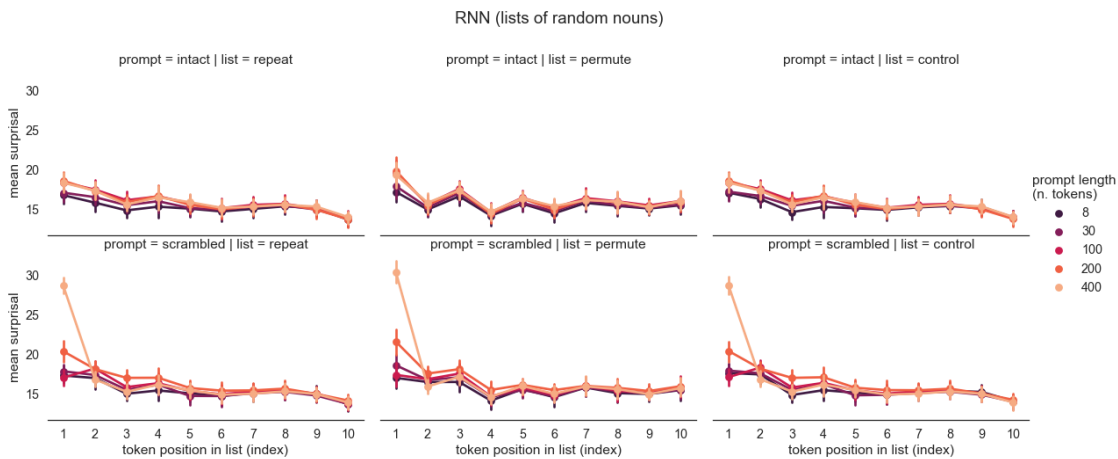
This suggests that increasing once we interfere with 400 permuted tokens, the RNN was not expecting a content noun after a colon. The jump is rather non-linear and it fades after one token or so.

```
[7]: # define a wrapper function
def make_plot(data_frame, list_type, title):

    sns.set_style("white")
    p = sns.catplot(data=data_frame, kind="point", x="marker_pos", y="surp",
        hue="prompt_len", col="second_list", row="context",
        estimator=np.mean, ci=95.0,
        col_order=["repeat", "permute", "control"],
        row_order=["intact", "scrambled"],
        palette="rocket")
    p.fig.set_size_inches(18,7)
    p.fig.subplots_adjust(top=0.85)
    p.fig.suptitle(title)
    p.set_axis_labels("token position in list (index)", "mean surprisal")
    p._legend.set_title("prompt length\n(n. tokens)")
    p.set_titles("prompt = {row_name} | list = {col_name}")
    p.despine(left=True);
```

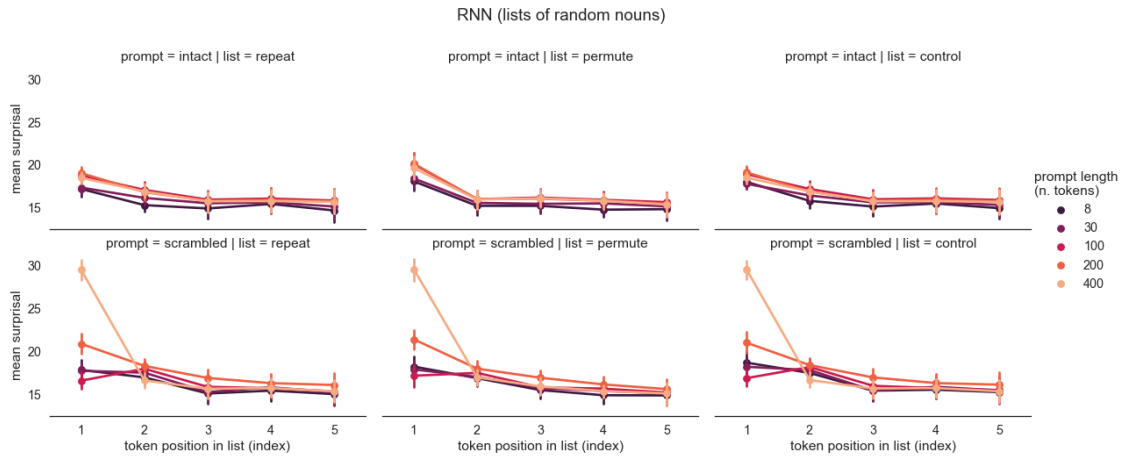
```
[8]: data=None
list_len=10
list_type = "random"
data=rnn.loc[(rnn.marker==3) & (rnn.list_len==list_len) & (rnn.
    list==list_type)] # list of len 10
```

```
[9]: make_plot(data, list_type, title="RNN (lists of {}) nouns".format(list_type))
```



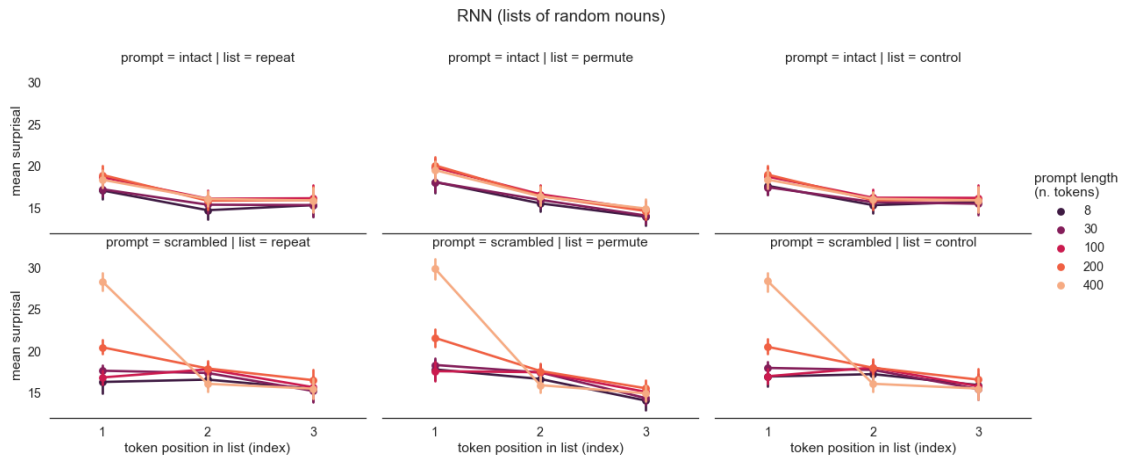
```
[10]: data=None
list_len=5
list_type = "random"
data=rnn.loc[(rnn.marker==3) & (rnn.list_len==list_len) & (rnn.
    list==list_type)] # list of len 10
```

```
[11]: make_plot(data, list_type, title="RNN (lists of {} nouns)".format(list_type))
```



```
[12]: data=None
list_len=3
list_type = "random"
data=rnn.loc[(rnn.marker==3) & (rnn.list_len==list_len) & (rnn.
↪list==list_type)] # list of len 10
```

```
[13]: make_plot(data, list_type, title="RNN (lists of {} nouns)".format(list_type))
```



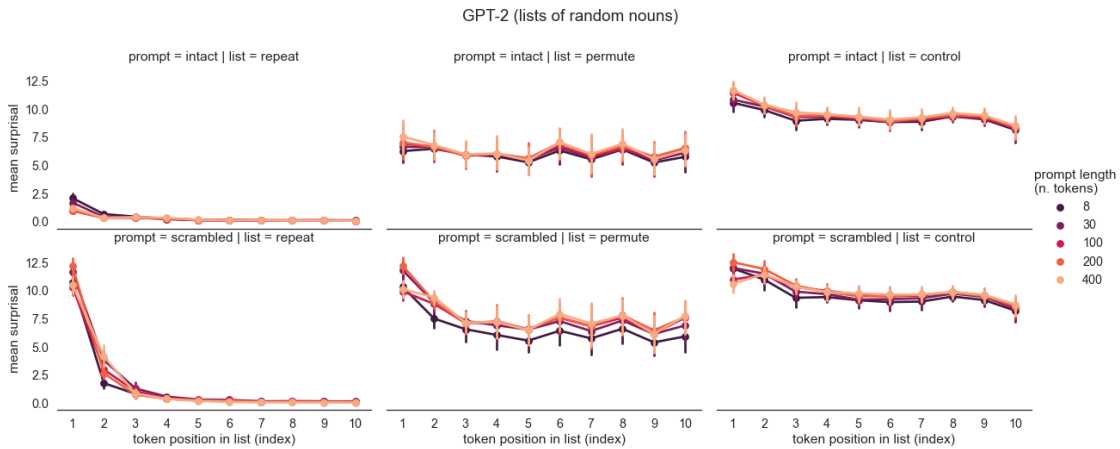
2.2 GPT-2

We repeat the same for the gpt-2 output.

As was clear before, GPT-2 is much more sensitive to list identity

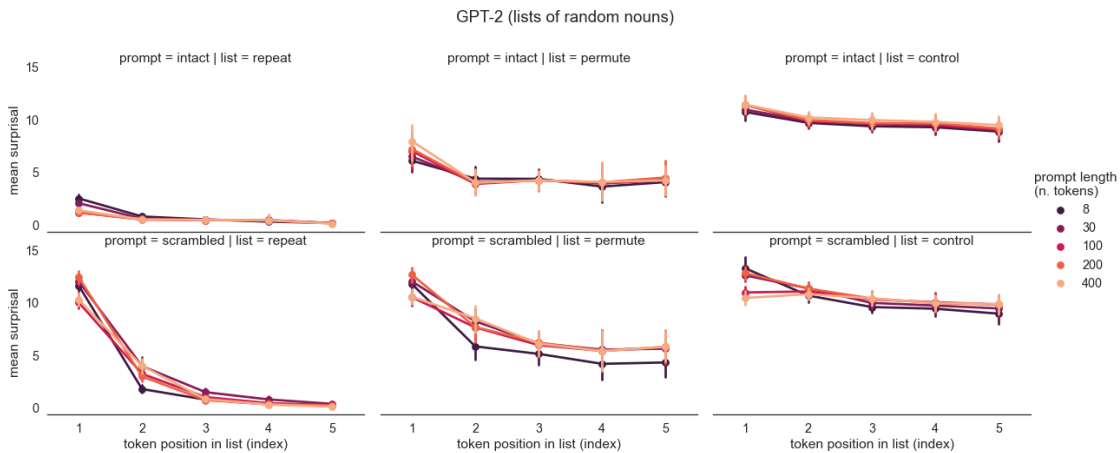
```
[14]: data=None
list_len=10
list_type = "random"
data=gpt.loc[(gpt.marker==3) & (gpt.list_len==list_len) & (gpt.list==list_type)
↳ & (~gpt.marker_pos.isin([11]))] # list of len 10
```

```
[15]: make_plot(data, list_type, title="GPT-2 (lists of {} nouns)".format(list_type))
```



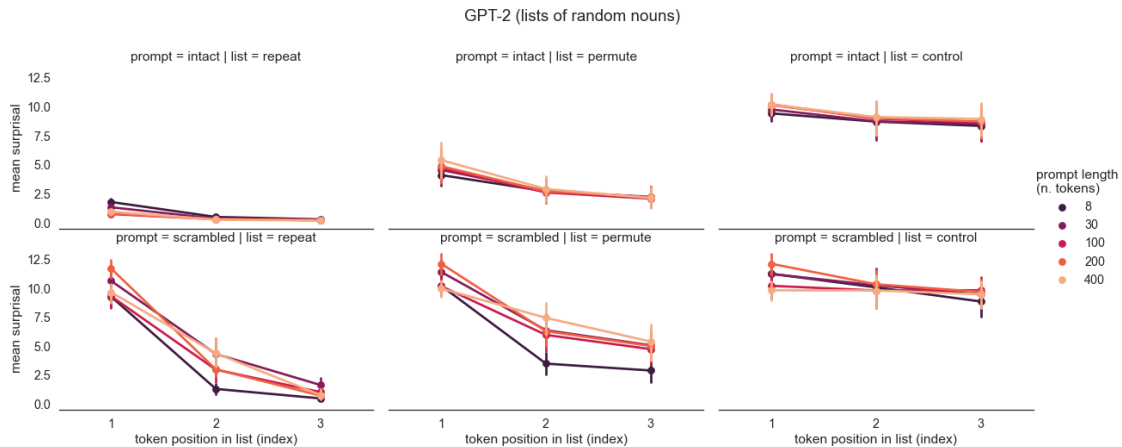
```
[16]: data=None
list_len=5
list_type = "random"
data=gpt.loc[(gpt.marker==3) & (gpt.list_len==list_len) & (gpt.list==list_type)
↳ & (gpt.marker_pos.isin(list(range(1, 6))))] # list of len 10
```

```
[17]: make_plot(data, list_type, title="GPT-2 (lists of {} nouns)".format(list_type))
```



```
[18]: data=None
list_len=3
list_type = "random"
data=gpt.loc[(gpt.marker==3) & (gpt.list_len==list_len) & (gpt.list==list_type)
↳ & (gpt.marker_pos.isin(list(range(1, 4))))] # list of len 10
```

```
[19]: make_plot(data, list_type, title="GPT-2 (lists of {} nouns)".format(list_type))
```



3 Interim conclusion

- GPT-2 stores list identity (flooring in list repeat condition) and word order (effect of repeating and permuting the list)
- cannot claim similar for RNN (either we cannot detect or there is no effect)
- if anything, GPT-2 seems more similar to a “classical working memory system” because, so far, shared semantic similarities (list repetition, prompt congruence) lead to facilitation (as inferred from reduced surprisal levels) on processing.
- if context (prompt) is incongruent and > 400 tokens length, the structure of context matters to RNN it seems, so it is not so idle after all
- so far, it is hard to say how number of items impacts the memory system